

1. Environment  
Overview

2. Installation and  
Verification

3. Deploying Container  
Native Storage

4. Application  
Management Basics

5. Project Template,  
Quota, and Limits

6. Using External  
Authentication  
Providers (LDAP)

7. Infrastructure  
Management Basics

8. Container-native  
Storage Management

9. Skipping modules

# Project Request Template, Quotas, Limits

5

In the Application Management Basics lab, you dealt with the fundamental building blocks of OpenShift, all contained within a **Project**.

Out of the box, OpenShift does not restrict the quantity of objects or amount of resources (eg: CPU, memory) that they can consume within a **Project**. Further, users may create an unlimited number of **Projects**. In a world with no constraints, or in a POC-type environment, that would be fine. But reality calls for a little restraint.

Before continuing, make sure you are using logged with the built-in system administrator account in OpenShift.



```
oc login -u system:admin
```

## Background: Project Request Template

When users use the `new-project` command, what is actually happening behind the scenes is a project request flow. Much like it sounds, there is a **Template** that is processed during the project request.

[View the Default Project Request Template](#)

Execute the following command to view the built-in default **Project Request Template**:

```
oc adm create-bootstrap-project-template
```

If you dig into the JSON output of this command, you will notice that there are various parameters defined at the end. Take a look at the help output for the **new-project** command:

```
oc new-project -h
```

Do you see how there is a **--display-name** directive on **new-project** and how there is a **PROJECT\_DISPLAYNAME** parameter?

The **new-app** workflow involves the user providing information to fulfill the project request. OpenShift decides if this request should be allowed (for example, are users allowed to create **Projects**? Does this user have too many **Projects**?) and, if so, processes the **Template**.

If you look at the objects defined in the **Template**, you will notice that there's no mention of quota or limits. You'll change that now.



**Templates** are a powerful tool that enable you to easily create reusable sets of OpenShift objects with powerful parameterization. They can be used to quickly deploy more complicated and related components into OpenShift, and can be a useful part of your organization's software development lifecycle (SDLC). More information can be found in the [template documentation](#). We will not go into more detail on **Templates** in these exercises.

# Modify the Project Request Template

You won't actually have to make template changes in this lab — we've made them for you already. Use `cat`, `less`, or your favorite editor to view the modified **Project Request Template**:

```
cat /opt/lab/support/project_request_template.yaml
```

Take note that there are two new sections added: **ResourceQuota** and **LimitRange**.

## Background: ResourceQuota

The [quota documentation](#) provides a great description of what **ResourceQuota** is about:

```
A resource quota, defined by a ResourceQuota object, provides constraint
limit aggregate resource consumption per project. It can limit the quant
objects that can be created in a project by type, as well as the total a
compute resources and storage that may be consumed by resources in that
project.
```

In our case, we are setting a specific set of quota for CPU, memory, storage, volume claims, and **Pods**. Take a look at the **ResourceQuota** section from the `project_request_template.yaml` file:

```
- apiVersion: v1
  kind: ResourceQuota
  metadata:
```

```
name: ${PROJECT_NAME}-quota ❶
spec:
  hard:
    pods: 10 ❷
    requests.cpu: 4000m ❸
    requests.memory: 8Gi ❹
    resourcequotas: 1
    requests.storage: 25Gi ❺
    persistentvolumeclaims: 5 ❻
```

- ❶ While only one quota can be defined in a **Project**, it still needs a unique name/id.
- ❷ The total number of pods in a non-terminal state that can exist in the project.
- ❸ CPUs are measured in "milicores", and more information on how Kubernetes/OpenShift calculates cores can be found in the [upstream documentation](#).
- ❹ There is a system of both **limits** and **requests** that we will discuss more when we get to the **LimitRange** object.
- ❺ Across all persistent volume claims in a project, the sum of storage requested cannot exceed this value.
- ❻ The total number of persistent volume claims in a project.

For more details on the available quota options, refer back to the [quota documentation](#).

# Background: LimitRange

The [limit range documentation](#) provides some good background:

A limit range, defined by a LimitRange object, enumerates compute resource constraints in a project at the pod, container, image, image stream, and persistent volume claim level, and specifies the amount of resources the container, image, image stream, or persistent volume claim can consume.

While the quota sets an upper bound on the total resource consumption within a project, the **LimitRange** generally applies to individual resources. For example, setting how much CPU an individual **Pod** or container can use.

Take a look at the sample **LimitRange** definition that we have provided in the `project_request_template.yaml` file:

```
- apiVersion: v1
  kind: LimitRange
  metadata:
    name: ${PROJECT_NAME}-limits
    creationTimestamp: null
  spec:
    limits:
    -
      type: Container
      max: 1
        cpu: 4000m
        memory: 1024Mi
      min: 2
        cpu: 10m
        memory: 5Mi
      default: 3
```

```
cpu: 4000m
memory: 1024Mi
defaultRequest: 4
cpu: 100m
memory: 512Mi
```

The difference between requests and default limits is important, and is covered in the [limit range documentation](#). But, generally speaking:

- 1 **max** is the highest value that may be specified for limits and requests
- 2 **min** is the lowest value that may be specified for limits and requests
- 3 **default** is the maximum amount (limit) that the container may consume, when nothing is specified
- 4 **defaultRequest** is the minimum amount that the container may consume, when nothing is specified

In addition to these topics, there are things like **Quality of Service Tiers** as well as a **Limit : Request** ratio. There is additionally more information in the [compute resources](#) section of the documentation.

For the sake of brevity, suffice it to say that there is a complex and powerful system of Quality of Service and resource management in OpenShift. Understanding the types of workloads that will be run in your cluster will be important to coming up with sensible values for all of these settings.

The settings we provide for you in these examples generally restrict projects to:

- A total CPU quota of 4 cores (4000m) where
  - Individual containers
    - must use 4 cores or less
    - cannot be defined with less than 10 milicores
    - will default to a request of 100 milicores (if not specified)
    - may burst up to a limit of 4 cores (if not specified)
- A total memory usage of 8 Gibibyte (8192 Megabytes) where
  - Individual containers
    - must use 1 Gi or less
    - cannot be defined with less than 5 Mi
    - will default to a request of 512 Mi
    - may burst up to a limit of 1024 Mi
- Total storage claims of 25 Gi or less
- A total number of 5 volume claims
- 10 or less **Pods**

In combination with quota, you can create very fine-grained controls, even across projects, for how users are allowed to request and utilize OpenShift's various resources.



Remember that quotas and limits are applied at the **Project** level. **Users** may have access to multiple **Projects**, but quotas and limits do not apply directly to **Users**. If you want to apply one quota

across multiple **Projects**, then you should look at the [multi-project quota](#) documentation. We will not cover multi-project quota in these exercises.

## Installing the Project Request Template

OK, with this background in place, let's go ahead and actually tell OpenShift to use this new **Project Request Template**.

### Create the Template

As we discussed earlier, a **Template** is just another type of OpenShift object. The `oc` command provides a `create` function that will take YAML/JSON as input and simply instantiate the objects provided.

Go ahead and execute the following:

```
oc create -f /opt/lab/support/project_request_template.yaml -n default
```

This will create the **Template** object in the `default` **Project**. You can now see the **Templates** in the `default` project with the following:

```
oc get template -n default
```

You will see something like the following:

NAME	DESCRIPTION	PARAMETERS	OBJECTS
project-request		5 (5 blank)	7



## Edit the `master-config.yaml`

For this exercise, we have not already configured things for you. Use your favourite editor (`vim`, `vi`, `nano`, etc.) to edit the master's configuration file using `sudo` privileges. For example:

```
sudo vim /etc/origin/master/master-config.yaml
```



If you are unfamiliar with the `vi` editor, you should probably use the `nano` editor. Use `^O` (Control + capital O) to save/write out the file after editing, and then `^X` to exit.

In the master's configuration file, you will find a line that mentions `projectRequestTemplate`. It will not have anything specified. When nothing is specified, OpenShift uses the built-in **Template** that you exported in the first exercise.

You will want to edit the config to look like the following (just that section):

```
...
projectConfig:
  projectRequestTemplate: "default/project-request"
...
```

## Restart the Master

Since you have made a configuration change to the master, you will need to restart its service. You can do so with the following command with `sudo` privileges:

```
sudo systemctl restart atomic-openshift-master
```

## Test the Project Request Template

At this point you have reconfigured the master to use the **Project Request Template** (a special kind of **Template**) called `project-request` that is located in the `default` **Project**. Now it is time to observe this change in action.

## Create a New Project

When creating a new project, you should see that a **Quota** and a **LimitRange** are created with it. First, create a new project called `template-test`:

```
oc new-project template-test
```

Then, use `describe` to look at some of this **Project's** details:

```
oc describe project template-test
```

The output will look something like:

```
Name:          template-test
Namespace:     <none>
Created:       5 seconds ago
Labels:        <none>
Annotations:   openshift.io/description=
               openshift.io/display-name=
               openshift.io/requester=fancyuser1
```

```

openshift.io/sa.scc.mcs=s0:c12,c9
openshift.io/sa.scc.supplemental-groups=1000150000/10000
openshift.io/sa.scc.uid-range=1000150000/10000
Display Name: <none>
Description: <none>
Status: Active
Node Selector: <none>
Quota:
  Name: template-test-quota
  Resource      Used    Hard
  -----
  persistentvolumeclaims 0      5
  pods                0      10
  requests.cpu         0      4
  requests.memory      0      4Gi
  requests.storage     0      25Gi
  resourcequotas       1      1
Resource limits:
  Name: template-test-limits
  Type   Resource      Min    Max    Default
  ----
  Container  cpu          10m    4      4
  Container  memory       5Mi    1Gi    1Gi

```

You can also see that the **Quota** and **LimitRange** objects were created:

```
oc get quota -n template-test
```

You will see:

```

NAME                AGE
template-test-quota 2m

```

And:

```
oc get limitrange -n template-test
```

You will see:

NAME	AGE
template-test-limits	2m

## Clean Up

If you wish, you can deploy the application from the Application Management Basics lab again inside this `template-test` project to observe how the **Quota** and **LimitRange** are applied. If you do, be sure to look at the JSON/YAML output (`oc get ... -o yaml`) for things like the **DeploymentConfig** and the **Pod**.

Before you continue, you may wish to delete the **Project** you just created:

```
oc delete project template-test
```

[Go to previous module](#)[Go to next module](#)

